


Exploring Caching Policies for Host-side Flash Caches

Bhavesh Singh
143059003

Guided by
Prof. Purushottam Kulkarni

Setting the Context - Centralized Storage

- Storage separate from servers
 - Accessible over the network
 - Exposed as logical disks
- Benefits
 - Storage can scale independently
 - Easy maintenance
- Challenges
 - Network round trips for each read/write.

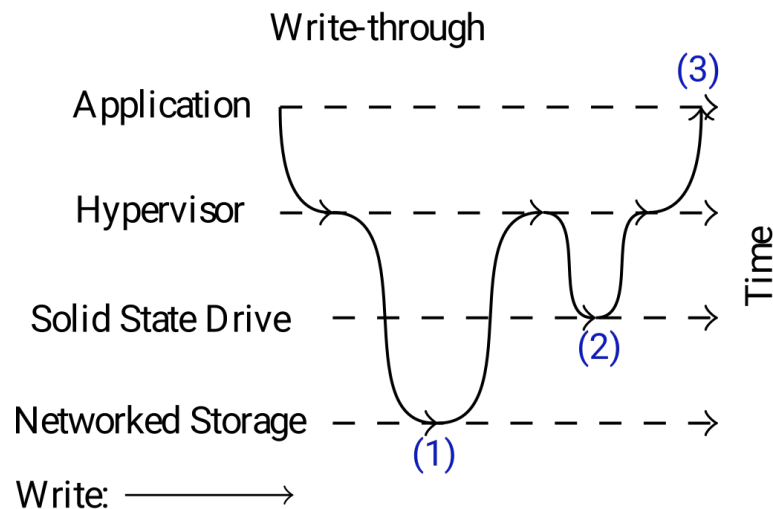


Setting the Context - Host-side Flash Caches

- Solid state disks used at the host-side as caching layer
- Benefits
 - Save network round trips to networked storage
 - Faster than network access
 - Persistent
- Challenges
 - Caching policy decisions
 - Data consistency


Caching Policies - Write Through

- Ack a write only after writing to storage and cache
- Good for read-heavy workloads
- Poor performance for write-heavy workloads
- Transactional Consistency
 - Each write acknowledged to the application is guaranteed to be on the storage




Caching Policies - Write Back

- Cache the writes too
- Ack a write after writing to cache
 - Evict on memory pressure
- Wonderful for write-heavy workloads
 - Write coalescing
- No consistency guarantees




Caching Policies - Write Back Variants

- Room for improvement
 - Trade off data staleness for improved write performance
 - Get lower consistency guarantees than write through
 - Batch updates to networked storage in the background



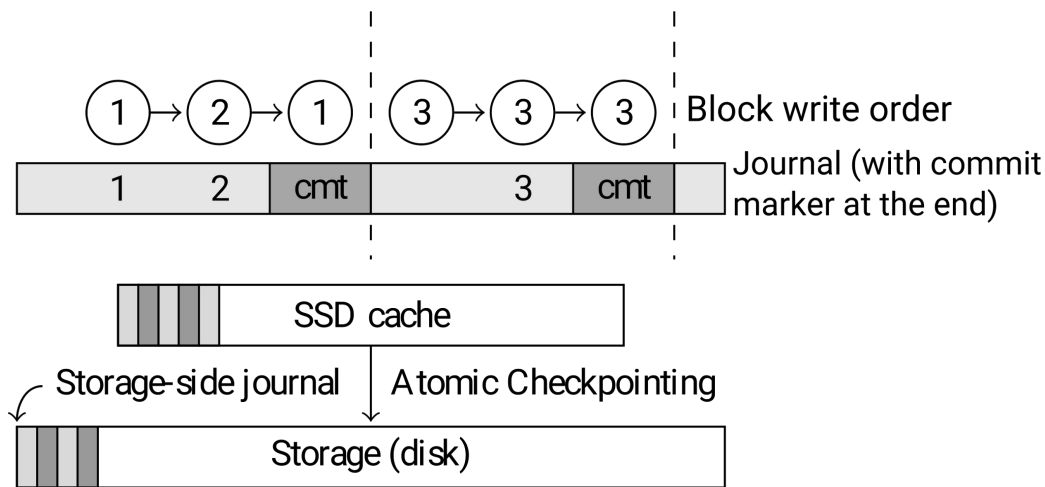
Ordered Write Back [Koller, FAST 13]

- Imposes a partial order on writes
 - Called completion-issue order
 - Ensure writes are evicted in order
- Point-in-time consistency
 - After recovery storage reflects application's view of storage at some point-in-time
- Space overhead
 - Each write to a block mapped to different cache block
- Data lost in case of flash failure




Journalled Write Back [Koller, FAST 13]

- Groups updates into journals
 - Journals committed atomically to networked storage
 - Writes within a journal can be coalesced
- Point-in-time consistency
- Data lost in case of flash failure



Write Back (with Peer Replication) [Bhagwat, FAST 15]


- Adds fault-tolerance
 - Each write synchronously replicated to flash caches of k peers
- Point-in-time consistent
- Can tolerate k flash failures and $k+1$ host failures
- Synchronous replication overhead



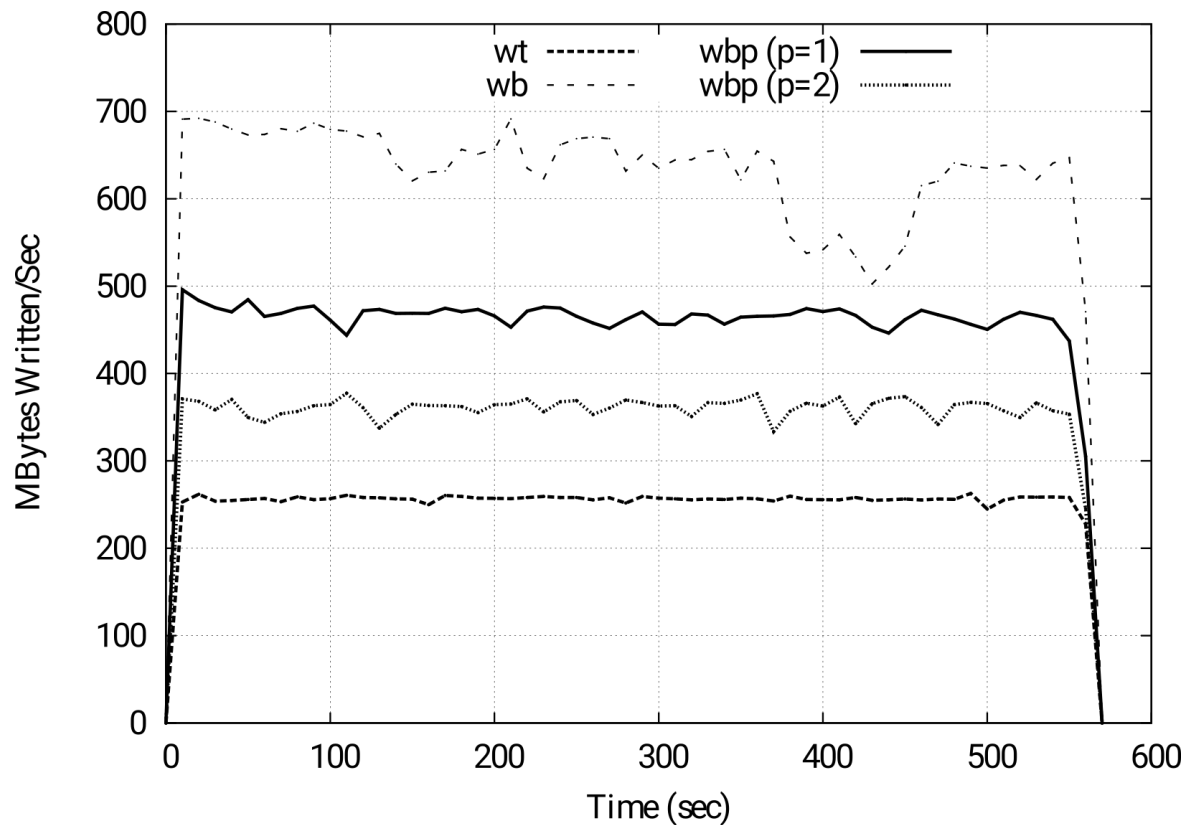
Write Back Flush and Write Back Persist [Qin, ATC 14]

- Flash cache provides consistency guarantees equivalent to the underlying storage
 - Applications only expect data to be persistent at write barriers viz. fsync.
- Application-level consistent
 - After recovery, storage reflects a state which satisfies some application-level constraints

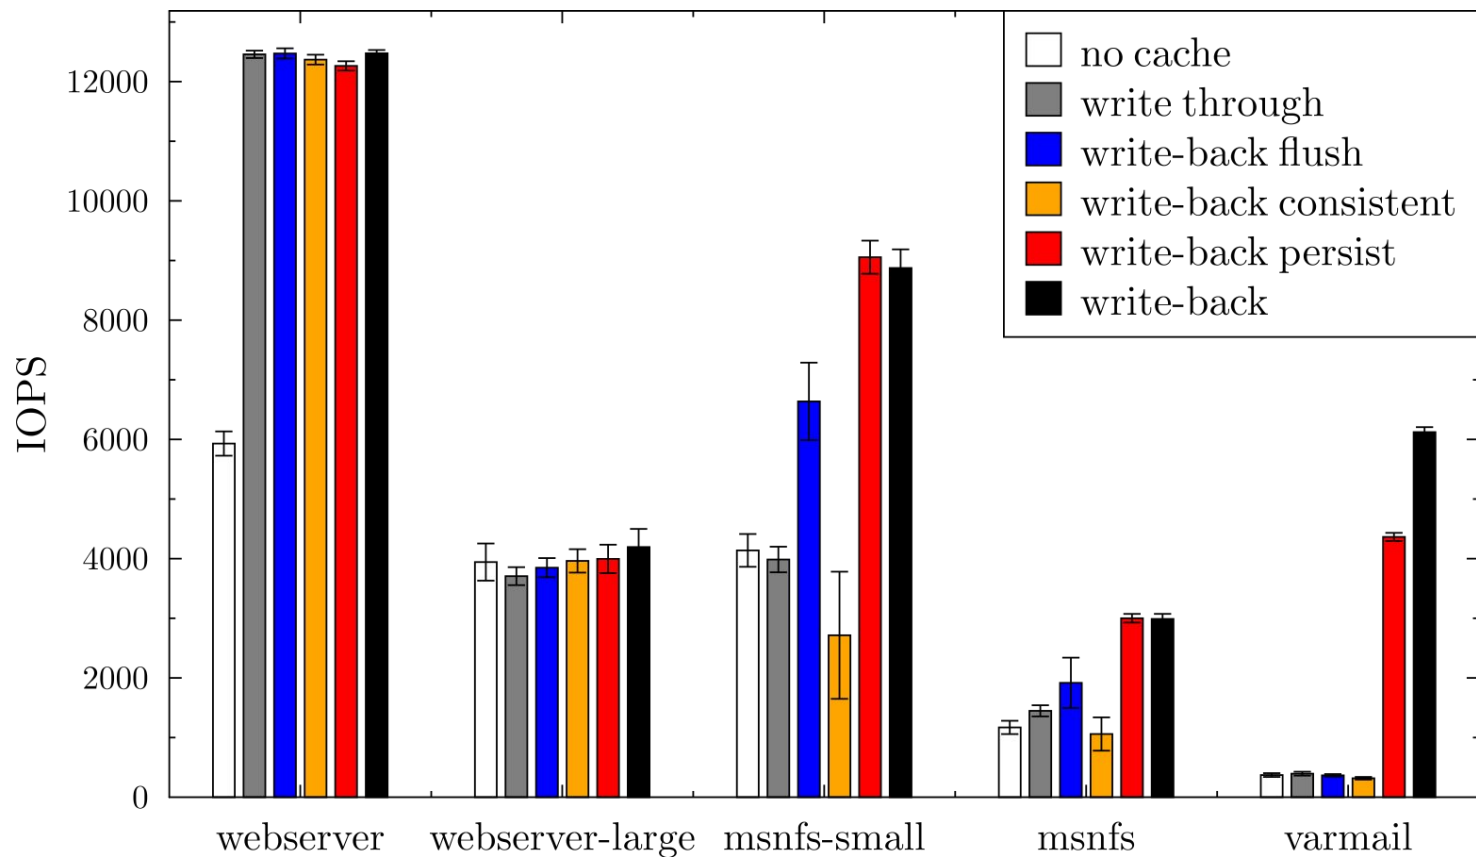
Performance Comparison [Koller, FAST 13]



Performance Comparison [Bhagwat, FAST 15]



Performance Comparison [Qin, ATC 14]



Overview of Caching Policies

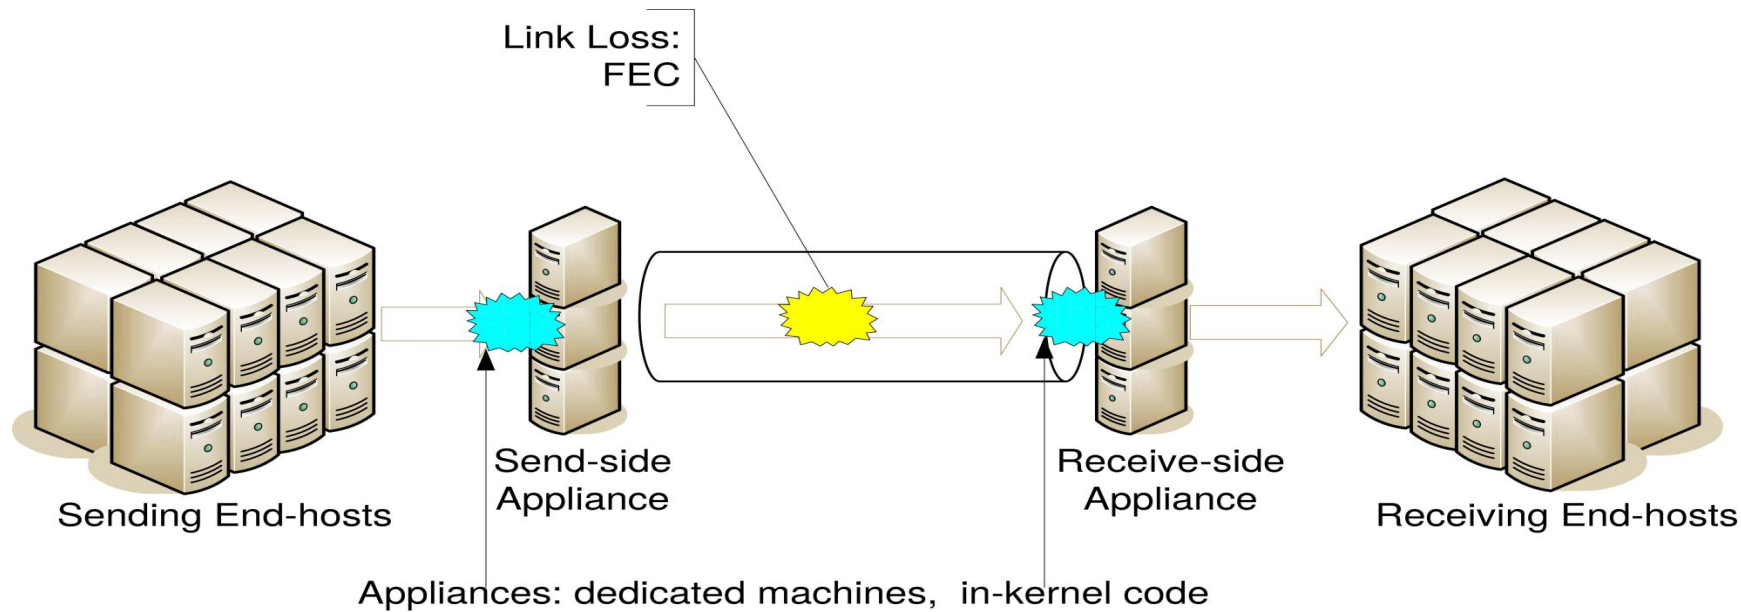
| Caching Policy | Consistency Provided | Fault Tolerance |
|--------------------------------|-----------------------------|------------------------|
| Write through | Transactional | |
| Write back | Point-in-time | |
| Ordered write back | Point-in-time | |
| Journalled write back | Point-in-time | |
| Write back flush | Application-level | |
| Write back persist | Application-level | Flash failures only |
| Write back with peering | Point-in-time | Yes |

Can We Do Better?

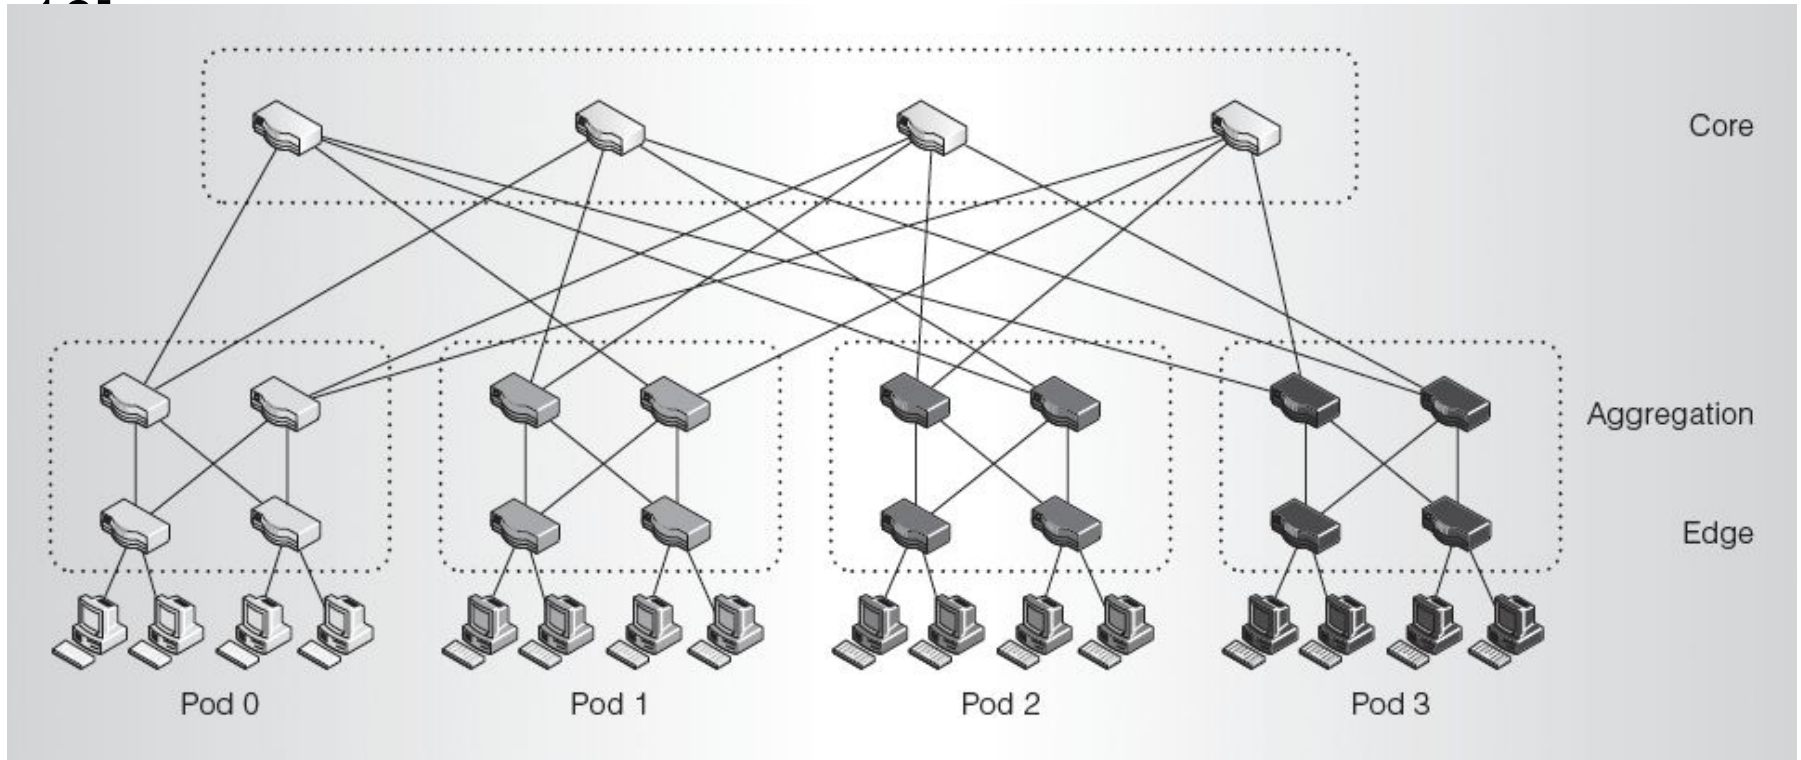
- GOAL: Be “almost synchronous” while performing fully asynchronously.
- Motivation
 - Using packet-level FEC [Balakrishnan, NSDI 08]
 - Monitor the network and create a global view to help choose which paths to send data on [Al-Fares, NSDI 10].

Packet-level FEC [Balakrishnan, NSDI 08]

Forward Error Correction used to recover from packet losses



Global View of Data Center Network [Al-Fares, NSDI




Our Solution - Probabilistic Asynchronous Writes

- Asynchronous

- Writes ack'd as soon as persisted on flash and put on the wire.

- Probabilistic

- Determine the loss probabilities of the packet on various routes to the networked storage and choose best path.
- Use packet-level FEC to overcome packet losses.



References

1. Koller, R., Marmol, L., Rangaswami, R., Sundararaman, S., Talagala, N., and Zhao, M. **Write Policies for Host-side Flash Caches.** *USENIX FAST 13.*
2. Bhagwat, D., Patil, M., Ostrowski, M., Vilayannur, M., Jung, W., and Kumar, C. **A Practical Implementation of Clustered Fault Tolerant Write Acceleration in a Virtualized Environment.** *USENIX FAST 15.*
3. Qin, D., Brown, A. D., and Goel, A. **Reliable Writeback for Client-side Flash Caches.** *USENIX ATC 14.*
4. Balakrishnan, M., Marian, T., Birman, K., Weatherspoon, H., and Vollset, E. **Maelstrom: Transparent Error Correction for Lambda Networks.** *USENIX NSDI 08.*
5. Al-Fares, M., Radhakrishnan, S., Raghavan, B., Huang, N., and Vahdat, A. **Hedera: Dynamic Flow Scheduling for Data Center Networks.** *USENIX NSDI 10.*

Thank You

And May the Fourth Be With You!
(Today is Star Wars Day)